

# SOFTENG 351: Fundamentals of Database Systems - Guide to Test 1

Michael Kemp

April 11, 2017

## Contents

<b>1</b>	<b>Basics of the relational model of data</b>	<b>2</b>
1.1	The relational model of data . . . . .	2
1.2	SQL as a data definition and query language . . . . .	2
<b>2</b>	<b>Query Languages</b>	<b>3</b>
2.1	Relational algebra . . . . .	3
2.2	Relational calculus . . . . .	4
2.3	SQL . . . . .	4
<b>3</b>	<b>Database design</b>	<b>4</b>
3.1	Entity-Relationship modelling . . . . .	4

# 1 Basics of the relational model of data

## 1.1 The relational model of data

- *relations* are sets of *tuples* often represented as a table
- *attributes* are the column titles of a relation
- for each attribute we assign a *domain* which is a universal set containing all possible values (like a string;  $\text{dom}(A) = \text{string}$ )
- *tuples* are the rows of a relation and all have the same structure in a relation
- if there is no value for an attribute then the value is *null*
- *relation schema* are a finite set  $R$  where attributes are  $A$  and each attribute  $A \in R$  has a domain  $\text{dom}(A)$
- relation schema can be written  $R = \{A_1, A_2 \dots A_n\}$  or  $R(A_1, A_2 \dots A_n)$  or  $R(A_1 : \text{dom}(A_1) \dots A_n : \text{dom}(A_n))$
- All  $R$ -*tuples* (a tuple in a relation schema) are an element  $t$  of the Cartesian product of the domains of all the attributes  $t \in A_1 \times A_2 \times \dots \times A_n$  because each attribute's value is bound to its respective domain.
- $R$ -*relations* are a finite set  $r$  of  $R$ -tuples thus  $r \subseteq \text{dom}(A_1) \times \dots \times \text{dom}(A_n)$
- $R$ -tuples can be written with their values  $t = (A_1 : v_1 \dots A_n : v_n)$
- A *database-schema* is a finite set  $S$  of relation schemata
- An  $S$ -database  $I$  consists of one  $R$ -relation for  $I(R)$  for each relation  $R$  in  $S$  ( $I = \{I(R) | R \in S\}$ )
- Having duplicates in a database is normally useless so we have *keys* to ensure a uniqueness over an attribute or a combination of
- *superkey* over a relation schema  $R$ 
  - finite, non-empty subset  $K \subseteq R$
  - is *satisfied* if an  $R$ -relation  $r$  only has tuples with a unique combination of values for each attribute in the superkey.
  - A *Key* is a superkey if there is no subset which is also a satisfied superkey
  - A *foreign key* is when all of the combination of values of attributes (in the foreign key) is in the set of the table which defines the foreign key ( $[A_1 \dots A_n] \subseteq S[A_1 \dots A_n]$ ). Also the same  $S$  can not be referenced twice.

## 1.2 SQL as a data definition and query language

- is a DDL (Data Definition Language) and a DML (Data Manipulation Language)
- names are not case sensitive
- DDL
  - CREATE TABLE  $\langle \text{tablename} \rangle \langle \text{attribute1 domain}, \dots \rangle$  attributes can be specified as NOT NULL
  - DROP TABLE  $\langle \text{tablename} \rangle$
  - ALTER TABLE followed by other stuff
- Domains
  - CHARACTER: Character strings of set length
  - VARCHAR: Character strings up to set length
  - NATIONAL CHARACTER
  - INTEGER: Signed 32-bit integer

- SMALLINT: Signed 16-bit integer
- NUMERIC = DECIMAL: Numeric values with definable precision and scale
- FLOAT: Approximate numeric values with definable precision (up to 64)
- REAL: Approximate numeric values up to 64 precision
- DOUBLE PRECISION: Double a REAL
- BIT = BOOLEAN: TRUE, FALSE, false or true
- BIT VARYING: Bit strings
- DATE: YYYY-MM-DD but can use single digits for month and date
- TIME: HH:MM:SS with optional nano seconds and seconds up to (including) 61.999999
- to insert a tuple to a relation all values must be known thus null exists if it doesn't exist or is not yet known
- *duplicate tuples* are tuples where both values are the same for every attribute
- duplicate tuples are not allowed in relations but IRL it's allowed because it's computationally expensive to manage.
- A table is  $X$ -total if all its tuples are  $X$ -total
- Constraints
  - NOT NULL means  $A$  must be  $A$ -total
  -

I'll come back to this

## 2 Query Languages

### 2.1 Relational algebra

- $A$  is the set of possible relations
- Partial operations on  $A$  take either 1 or 2 relations as input and produce another relation as output
- $\#r$  is the set of attributes of a relation  $r$
- Operations:
  - Attribute selection:  $\sigma_{A=B}$  produces a relation where tuples have the same value in attribute  $A$  as in  $B$
  - Constant selection:  $\sigma_{A=c}$  Same as attribute but against a constant  $c$  instead of an attribute
  - Projection:  $\pi_{A_1, \dots, A_n}(r)$  Takes a relation  $r$  and returns another with only the specified attributes  $A_1, \dots, A_n$
  - Renaming:  $\delta_{oldname \rightarrow newname}$  Changes the name of an attribute without changing the relation
  - Union:  $r_1 \cup r_2$  Relations with the same set of attributes  $\#r_1 = \#r_2$  form a relation with tuples from both
  - Difference:  $r_1 - r_2$  Relations with the same set of attributes  $\#r_1 = \#r_2$  form a relation with the tuples from the first relation  $r_1$  that aren't in the second relation  $r_2$
  - Join (Natural):  $r_1 \bowtie r_2$  Joins tuples in both relations where attributes that are in both all have the same values. If there is more than one match then the cross product is given
- Redundant Operations:
  - Intersection:  $r_1 \cap r_2$  Returns a relation where tuples are in both relations  $r_1, r_2$  and have the same set of attributes  $\#r_1 = \#r_2$
  - Cross-product:  $r_1 \times r_2$  where both relation have no common attributes the produced relation is every possible pair of tuples
  - Division:  $r_1 \div r_2$  the set of attributes in  $r_2$  must be a superset of  $r_1$ . Where there is there is attribute values for every tuple in  $r_2$ . It's hard to explain and much easier to show. <http://www.mathcs.emory.edu/~cheung/Courses/377/Syllabus/4-RelAlg/division.html>

## **2.2 Relational calculus**

tell the author to stop procrastinating...

## **2.3 SQL**

tell the author to stop procrastinating...

# **3 Database design**

## **3.1 Entity-Relationship modelling**

tell the author to stop procrastinating...